

Deep learning in neuroscience

Tutorial 1

Carsen Stringer, PhD




About me



HHMI Janelia

How do neurons perform complex, high-dimensional computations?




Thanks to tutorial c

And thanks to tutorial

Roosbeh Farhoodi

Madineh Sarvestani

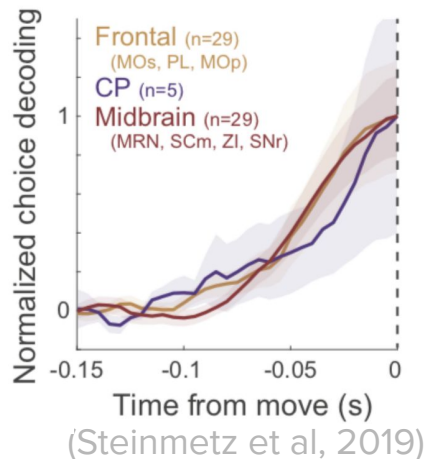
Ella Batty



Decoding models vs encoding models

(W1D4 tutorial 2)
neural activity \Rightarrow external correlate

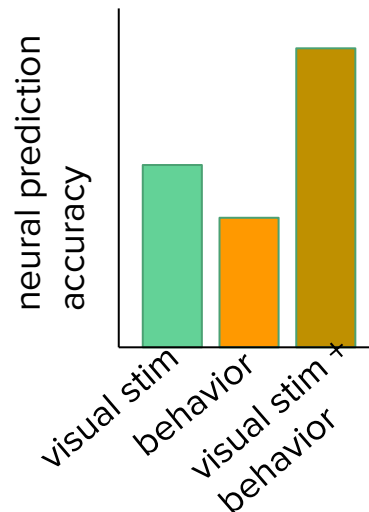
TUTORIAL 1



How much information does brain area X have about external correlate Y?

(W1D4 tutorial 1)
external correlate \Rightarrow neural activity


TUTORIAL 2



How does brain area X represent external correlate Y?



Neural recordings in mice



resp contains responses of 23589 neurons to 360 stimuli


Stringer, Michaelos, Pachitariu, *bioRxiv*, 2019



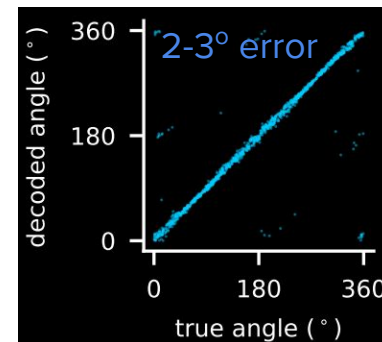
Building a deep network

to predict *stimuli* from *neural responses*


Building a linear network




Can we predict better using a deep network?



Building a linear network



Creating a linear network in pytorch



run the network

```
import torch
from torch import nn

class DeepNet(nn.Module):
    def __init__(self, n_inputs, n_hidden):
        super().__init__() # needed to invoke the properties of the parent class nn.Module
        self.in_layer = nn.Linear(n_inputs, n_hidden) # neural activity --> hidden units
        self.out_layer = nn.Linear(n_hidden, 1) # hidden units --> output

    def forward(self, r):
        h = self.in_layer(r) # hidden representation
        y = self.out_layer(h)
        return y

# Initialize a deep network with M=200 hidden units
net = DeepNet(n_neurons, 200)

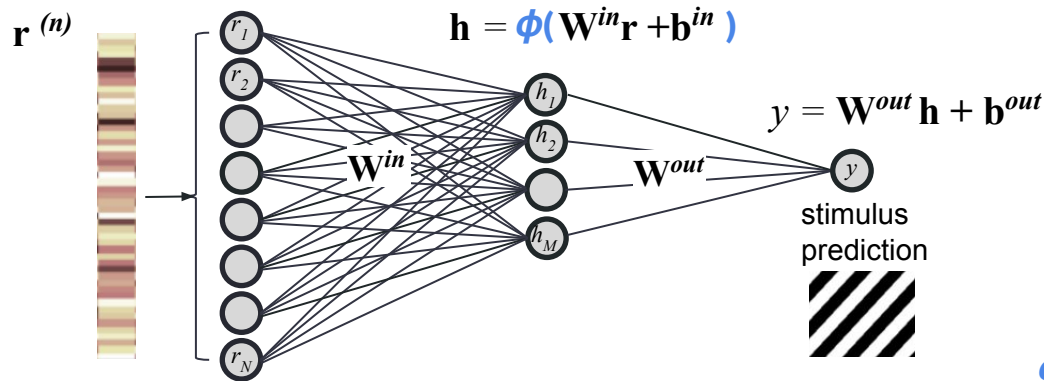
# Get neural responses to first stimulus in the data set
istim = 0 # index of first stimulus
r = resp[istim] # neural responses to this stimulus
# Decode orientation from these neural responses using initialized network
out = net(r) # compute output from network, equivalent to net.forward(r)
```

Activation functions



Activation functions

add non-linearities that allow flexible fitting




$$\frac{\partial \phi}{\partial \mathbf{x}} = \begin{cases} 1, & \text{if } \mathbf{x} \geq 0 \\ 0, & \text{if } \mathbf{x} < 0 \end{cases}$$

RELU


$$\phi(x) = \max(0, x)$$

$\phi(x)$



RELU activation function in pytorch

```
class DeepNetReLU(nn.Module):  
  
    def __init__(self, n_inputs, n_hidden):  
        super().__init__()  
        self.in_layer = nn.Linear(n_inputs, n_hidden)  
        self.out_layer = nn.Linear(n_hidden, 1)  
  
    def forward(self, r):  
        h = torch.relu(self.in_layer(r))  
        y = self.out_layer(h)  
        return y
```



Optimizing Neural Networks



Loss function

```
# Decode orientation from these neural responses
out = net(r) # compute output from network

# Initialize PyTorch mean squared error loss
function
loss_fn = nn.MSELoss()

# Evaluate mean squared error
loss = loss_fn(out, ori)
print('mean squared error: %.2f' % loss)
```

mean squared error: 12.73



Gradient descent

compute gradients wrt loss function L

$$\frac{\partial L}{\partial \mathbf{W}^{in}}, \frac{\partial L}{\partial \mathbf{b}^{in}}, \frac{\partial L}{\partial \mathbf{W}^{out}}, \frac{\partial L}{\partial \mathbf{b}^{out}}$$



update parameters by gradients on each iteration

$$\mathbf{W}^{in} \leftarrow \mathbf{W}^{in} - \alpha \frac{\partial L}{\partial \mathbf{W}^{in}}$$

$$\mathbf{b}^{in} \leftarrow \mathbf{b}^{in} - \alpha \frac{\partial L}{\partial \mathbf{b}^{in}}$$


$$\mathbf{W}^{out} \leftarrow \mathbf{W}^{out} - \alpha \frac{\partial L}{\partial \mathbf{W}^{out}}$$

$$\mathbf{b}^{out} \leftarrow \mathbf{b}^{out} - \alpha \frac{\partial L}{\partial \mathbf{b}^{out}}$$




Ella Batty

Loss function



adding skip connections



Li , Xu, Taylor, Studer, Goldstein, *arXiv*, 2018

Gradient descent in pytorch

```
optimizer = optim.SGD(net.parameters(), lr=.001)

for i in range(n_iter):
    # Evaluate loss
    out = net(train_data)
    loss = loss_fn(out, train_labels)

    # Compute gradients
    optimizer.zero_grad() # clear gradients
    loss.backward()

    # Update weights
    optimizer.step()
```



automatic differentiation!

you can use this for any complex model
you want to fit!



Gradient descent in pytorch

Did we need a deep network?



Regularization




Regularizing deep nets

Deep networks can have millions of parameters

How to prevent overfitting? \Rightarrow Regularization

- L2 regularization (weight decay)
- L1 regularization
- Batch normalization*
- Dropout
- Fewer neurons, fewer layers, use convolutions



Canziani et al, 2016



L2 regularization (weight decay)

$$\text{Cost} = ||Y_{target} - \text{net}(X_{input})||^2 + \lambda \sum ||W||^2$$

$$\partial \text{Cost} / \partial W = \partial / \partial W ||Y_{target} - \text{net}(X_{input})||^2 + 2\lambda \sum W$$

$$W \leftarrow W - \eta \cdot (\dots + 2\lambda W)$$

learning
rate

weight
decay

```
L2 = L2_penalty * torch.square(weights).sum( )
```

L1 regularization

L1 norm = abs


$$\text{Cost} = ||Y_{target} - \text{net}(X_{input})||^2 + \lambda \sum |W|^1$$

$$W \leftarrow W - \eta \cdot (\dots + \lambda \cdot \text{sign}(W))$$


- L1 produces sparse solutions
- L1 is more interpretable
- L1 is harder to optimize

```
L1 = L1_penalty * torch.abs(weights).sum()
```

A L1 regularization



B L2 regularization



C.V.Krishnakumar

